

**15-440**


# **MapReduce Programming**

## **Oct 30, 2012**

### **Topics**

- **Large-scale computing**
  - Traditional high-performance computing (HPC)
  - Cluster computing
- **MapReduce**
  - Definition
  - Examples
- **Implementation**
- **Alternatives to MapReduce**
- **Properties**

# Typical HPC Machine



## Compute Nodes

- High end processor(s)
- Lots of RAM

## Network

- Specialized
- Very high performance

## Storage Server

- RAID-based disk array

# HPC Machine Example


## Jaguar Supercomputer

- 3<sup>rd</sup> fastest in world



## Compute Nodes

- 18,688 nodes in largest partition
- 2X 2.6Ghz 6-core AMD Opteron
- 16GB memory
- Total: 2.3 petaflop / 300 TB memory




## Network

- 3D torus
  - Each node connected to 6 neighbors via 6.0 GB/s links

## Storage Server

- 10PB RAID-based disk array

# HPC Programming Model



- **Programs described at very low level**
  - Specify detailed control of processing & communications
- **Rely on small number of software packages**
  - Written by specialists
  - Limits classes of problems & solution methods

# Bulk Synchronous Programming

## Solving Problem Over Grid


- E.g., finite-element computation

## Partition into Regions


- $p$  regions for  $p$  processors

## Map Region per Processor

- Local computation sequential
- Periodically communicate boundary values with neighbors



# Typical HPC Operation



## Characteristics

- Long-lived processes
- Make use of spatial locality
- Hold all program data in memory (no disk access)
- High bandwidth communication


## Strengths

- High utilization of resources
- Effective for many scientific applications

## Weaknesses

- Requires careful tuning of application to resources
- Intolerant of any variability

# HPC Fault Tolerance



## Checkpoint

- Periodically store state of all processes
- Significant I/O traffic

## Restore

- When failure occurs
- Reset state to that of last checkpoint
- All intervening computation wasted

## Performance Scaling


- Very sensitive to number of failing components

# Google Data Centers




## Dalles, Oregon

- Hydroelectric power @ 2¢ / KW Hr
- 50 Megawatts
- Enough to power 60,000 homes



- Engineered for maximum modularity & power efficiency
- Container: 1160 servers, 250KW
- Server: 2 disks, 2 processors

# Typical Cluster Machine



## Compute + Storage Nodes

- Medium-performance processors
- Modest memory
- 1-2 disks

## Network

- Conventional Ethernet switches
  - 10 Gb/s within rack
  - 100 Gb/s across racks

# Machines with Disks

Lots of storage for cheap

- Seagate Barracuda
- 3 TB @ \$130  
(4.3¢ / GB)
- Compare 2007:  
0.75 TB @ \$266  
35¢ / GB



Seagate Barracuda 7200 3 TB 7200RPM SATA 6 Gb/s NCQ 64MB Cache 3.5-Inch Internal Bare Drive ST3000DM001

by [Seagate](#)

(378 customer reviews) | #1 Best Seller in Internal Hard Drives

List Price: \$269.99

Price: **\$129.99** & this item ships for **FREE with Super Saver Shipping.** [Details](#)

You Save: \$140.00 (52%)

**In Stock.**

Ships from and sold by [Amazon.com](#) in certified [Frustration-Free Packaging](#). Gift-wrap available.

## Drawbacks

- Long and highly variable delays
- Not very reliable

Not included in HPC Nodes

# Oceans of Data, Skinny Pipes



No more blaming connection speeds for your losses.

Verizon FiOS – the fastest Internet available.

Plans as low **\$39.99/month** (up to 5 Mbps).  
Plus, order online & **get your first month FREE!**

Enter your home phone number below to check availability.

**GO!**

Don't have a Verizon phone number? [Qualify your address.](#)



## 1 Terabyte

- Easy to store
- Hard to move

Disks	MB / s	Time
Seagate Barracuda	115	2.3 hours
Seagate Cheetah	125	2.2 hours
Networks	MB / s	Time
Home Internet	< 0.625	> 18.5 days
Gigabit Ethernet	< 125	> 2.2 hours
PSC Teragrid Connection	< 3,750	> 4.4 minutes


# Data-Intensive System Challenge

For Computation That Accesses 1 TB in 5 minutes


- Data distributed over 100+ disks
  - Assuming uniform data partitioning
- Compute using 100+ processors
- Connected by gigabit Ethernet (or equivalent)

## System Requirements

- Lots of disks
- Lots of processors
- Located in close proximity
  - Within reach of fast, local-area network




# Ideal Cluster Programming Model




- Application programs written in terms of high-level operations on data
- Runtime system controls scheduling, load balancing, ...

# Map/Reduce Programming Model



- **Map computation across many objects**
  - E.g.,  $10^{10}$  Internet web pages
- **Aggregate results in many different ways**
- **System deals with issues of resource allocation & reliability**

# MapReduce Example



Come, Dick.

Come and see.

Come, come.

Come and see.

Come and see Spot.

Come,  
Dick

Come  
and  
see.


Come,  
come.

Come  
and  
see.

Come  
and  
see  
Spot.

- Create an word index of set of documents

# MapReduce Example




- Map: generate  $\langle \text{word}, \text{count} \rangle$  pairs for all words in document
- Reduce: sum word counts across documents

# Hadoop Project



## File system with files distributed across nodes



- Store multiple (typically 3 copies of each file)
  - If one node fails, data still available
- Logically, any node has access to any file
  - May need to fetch across network

## Map / Reduce programming environment

- Software manages execution of tasks on nodes

# Hadoop MapReduce API

## Requirements

- Programmer must supply Mapper & Reducer classes

## Mapper

- Steps through file one line at a time
- Code generates sequence of <key, value> pairs
  - Call output.collect(key, value)
- Default types for keys & values are strings
  - Lots of low-level machinery to convert to & from other data types
  - But can use anything “writable”

## Reducer

- Given key + iterator that generates sequence of values
- Generate one or more <key, value> pairs
  - Call output.collect(key, value)

# Hadoop Word Count Mapper

```
public class WordCountMapper extends MapReduceBase
    implements Mapper {

    private final static Text word = new Text();

    private final static IntWritable count = new IntWritable(1);

    public void map(WritableComparable key, Writable values,
                    OutputCollector output, Reporter reporter)
        throws IOException {
        /* Get line from file */
        String line = values.toString();
        /* Split into tokens */
        StringTokenizer itr = new StringTokenizer(line.toLowerCase(),
                                                "\t.!:()[],'&-;|0123456789");
        while(itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            /* Emit <token,1> as key + value */
            output.collect(word, count);
        }
    }
}
```

# Hadoop Word Count Reducer

```
public class WordCountReducer extends MapReduceBase
    implements Reducer {

    public void reduce(WritableComparable key, Iterator values,
                       OutputCollector output, Reporter reporter)
        throws IOException {
        int cnt = 0;
        while(values.hasNext()) {
            IntWritable ival = (IntWritable) values.next();
            cnt += ival.get();
        }
        output.collect(key, new IntWritable(cnt));
    }
}
```

# Cluster Scalability Advantages

- Distributed system design principles lead to scalable design
- Dynamically scheduled tasks with state held in replicated files

## Provisioning Advantages

- Can use consumer-grade components
  - maximizes cost-peformance
- Can have heterogenous nodes
  - More efficient technology refresh

## Operational Advantages

- Minimal staffing
- No downtime

# Example: Sparse Matrices with Map/Reduce


$$\begin{matrix} \mathbf{A} \\ \left[ \begin{array}{ccc} 10 & 20 & \\ & 30 & 40 \\ 50 & 60 & 70 \end{array} \right] \end{matrix} \times \begin{matrix} \mathbf{B} \\ \left[ \begin{array}{ccc} -1 & & \\ -2 & -3 & \\ & -4 & \end{array} \right] \end{matrix} = \begin{matrix} \mathbf{C} \\ \left[ \begin{array}{ccc} -10 & -80 & \\ -60 & -250 & \\ -170 & -460 & \end{array} \right] \end{matrix}$$

- Task: Compute product  $\mathbf{C} = \mathbf{A} \cdot \mathbf{B}$
- Assume most matrix entries are 0

## Motivation


- Core problem in scientific computing
- Challenging for parallel execution
- Demonstrate expressiveness of Map/Reduce

# Computing Sparse Matrix Product



- Represent matrix as list of nonzero entries  
 $\langle \text{row}, \text{col}, \text{value}, \text{matrixID} \rangle$
- Strategy
  - Phase 1: Compute all products  $a_{i,k} \cdot b_{k,j}$
  - Phase 2: Sum products for each entry  $i,j$
  - Each phase involves a Map/Reduce

# Phase 1 Map of Matrix Multiply



- Group values  $a_{i,k}$  and  $b_{k,j}$  according to key k

# Phase 1 “Reduce” of Matrix Multiply

**Key = 1**

$$\begin{array}{r} 1 \xrightarrow[A]{10} 1 \\ 3 \xrightarrow[A]{50} 1 \end{array} \times \begin{array}{r} 1 \xrightarrow[B]{-1} 1 \end{array}$$

**Key = 2**

$$\begin{array}{r} 2 \xrightarrow[A]{30} 2 \\ 3 \xrightarrow[A]{60} 2 \end{array} \times \begin{array}{r} 2 \xrightarrow[B]{-2} 1 \\ 2 \xrightarrow[B]{-3} 2 \end{array}$$

**Key = 3**

$$\begin{array}{r} 1 \xrightarrow[A]{20} 3 \\ 2 \xrightarrow[A]{40} 3 \\ 3 \xrightarrow[A]{70} 3 \end{array} \times \begin{array}{r} 3 \xrightarrow[B]{-4} 2 \end{array}$$


$$\begin{array}{r} 1 \xrightarrow[C]{-10} 1 \\ 3 \xrightarrow[A]{-50} 1 \end{array}$$

$$\begin{array}{r} 2 \xrightarrow[C]{-60} 1 \\ 2 \xrightarrow[C]{-90} 2 \\ 3 \xrightarrow[C]{-120} 1 \\ 3 \xrightarrow[C]{-180} 2 \end{array}$$


$$\begin{array}{r} 1 \xrightarrow[C]{-80} 2 \\ 2 \xrightarrow[C]{-160} 2 \\ 3 \xrightarrow[C]{-280} 2 \end{array}$$

- Generate all products  $a_{i,k} \cdot b_{k,j}$

# Phase 2 Map of Matrix Multiply




Key = row,col



- Group products  $a_{i,k} \cdot b_{k,j}$  with matching values of i and j

# Phase 2 Reduce of Matrix Multiply



$$C = \begin{bmatrix} -10 & -80 \\ -60 & -250 \\ -170 & -460 \end{bmatrix}$$

- Sum products to get final entries

# Matrix Multiply Phase 1 Mapper

```
public class P1Mapper extends MapReduceBase implements Mapper {  
  
    public void map(WritableComparable key, Writable values,  
                    OutputCollector output, Reporter reporter) throws  
IOException {  
    try {  
        GraphEdge e = new GraphEdge(values.toString());  
        IntWritable k;  
        if (e.tag.equals("A"))  
            k = new IntWritable(e.toNode);  
        else  
            k = new IntWritable(e.fromNode);  
        output.collect(k, new Text(e.toString()));  
    } catch (BadGraphException e) {}  
}  
}
```

# Matrix Multiply Phase 1 Reducer

```
public class P1Reducer extends MapReduceBase implements Reducer {  
  
    public void reduce(WritableComparable key, Iterator values,  
                      OutputCollector output, Reporter reporter)  
        throws IOException  
{  
    Text outv = new Text(""); // Don't really need output values  
    /* First split edges into A and B categories */  
    LinkedList<GraphEdge> alist = new LinkedList<GraphEdge>();  
    LinkedList<GraphEdge> blist = new LinkedList<GraphEdge>();  
    while(values.hasNext()) {  
        try {  
            GraphEdge e =  
                new GraphEdge(values.next().toString());  
            if (e.tag.equals("A")) {  
                alist.add(e);  
            } else {  
                blist.add(e);  
            }  
        } catch (BadGraphException e) {}  
    }  
    // Continued
```

# MM Phase 1 Reducer (cont.)

```
// Continuation

Iterator<GraphEdge> aset = alist.iterator();
// For each incoming edge
while(aset.hasNext()) {
    GraphEdge aedge = aset.next();
    // For each outgoing edge
    Iterator<GraphEdge> bset = blist.iterator();
    while (bset.hasNext()) {
        GraphEdge bedge = bset.next();
        GraphEdge newe = aedge.contractProd(bedge);
        // Null would indicate invalid contraction
        if (newe != null) {
            Text outk = new Text(newe.toString());
            output.collect(outk, outv);
        }
    }
}
```

# Matrix Multiply Phase 2 Mapper

```
public class P2Mapper extends MapReduceBase implements Mapper {  
  
    public void map(WritableComparable key, Writable values,  
                    OutputCollector output, Reporter reporter)  
        throws IOException {  
        String es = values.toString();  
        try {  
            GraphEdge e = new GraphEdge(es);  
            // Key based on head & tail nodes  
            String ks = e.fromNode + " " + e.toNode;  
            output.collect(new Text(ks), new Text(e.toString()));  
        } catch (BadGraphException e) {}  
    }  
}
```

# Matrix Multiply Phase 2 Reducer

# Lessons from Sparse Matrix Example

## Associative Matching is Powerful Communication Primitive

- Intermediate step in Map/Reduce

## Similar Strategy Applies to Other Problems

- Shortest path in graph
- Database join

## Many Performance Considerations

- Kiefer, Volk, Lehner, TU Dresden
- Should do systematic comparison to other sparse matrix implementations

# MapReduce Implementation

## Built on Top of Parallel File System

- Google: GFS, Hadoop: HDFS
- Provides global naming
- Reliability via replication (typically 3 copies)


## Breaks work into tasks

- Master schedules tasks on workers dynamically
- Typically #tasks >> #processors


## Net Effect

- Input: Set of files in reliable file system
- Output: Set of files in reliable file system

# MapReduce Execution



# Mapping




## Hash Function $h$

- Maps each key  $K$  to integer  $i$  such that  $0 \leq i < R$


## Mapper Operation

- Reads input file blocks
- Generates pairs  $\langle K, V \rangle$
- Writes to local file  $h(K)$



# Mapping


- Dynamically map input file blocks onto mappers
- Each generates key/value pairs from its blocks
- Each writes R files on local file system



# Shuffling

## Each Reducer:


- Handles  $1/R$  of the possible key values
- Fetches its file from each of  $M$  mappers
- Sorts all of its entries to group values by keys



# Reducing


## Each Reducer:

- Executes reducer function for each key
- Writes output values to parallel file system



# MapReduce Effect

R Output Files




## MapReduce Step

- Reads set of files from file system
- Generates new set of files

Can iterate to do more complex processing

# Map/Reduce Operation

## Map/Reduce



## Characteristics

- Computation broken into many, short-lived tasks
  - Mapping, reducing
- Use disk storage to hold intermediate results

## Strengths

- Great flexibility in placement, scheduling, and load balancing
- Can access large data sets

## Weaknesses

- Higher overhead
- Lower raw performance

# Example Parameters

## Sort Benchmark

- **$10^{10}$  100-byte records**
- **Partition into  $M = 15,000$  64MB pieces**
  - Key = value
  - Partition according to most significant bytes
- **Sort locally with  $R = 4,000$  reducers**

## Machine

- **1800 2Ghz Xeons**
- **Each with 2 160GB IDE disks**
- **Gigabit ethernet**
- **891 seconds total**

# Interesting Features


## Fault Tolerance

- Assume reliable file system
- Detect failed worker
  - Heartbeat mechanism
- Reschedule failed task

## Stragglers

- Tasks that take long time to execute
- Might be bug, flaky hardware, or poor partitioning
- When done with most tasks, reschedule any remaining executing tasks
  - Keep track of redundant executions
  - Significantly reduces overall run time

# Map/Reduce Fault Tolerance




## Data Integrity

- Store multiple copies of each file
- Including intermediate results of each Map / Reduce
  - Continuous checkpointing

## Recovering from Failure

- Simply recompute lost result
  - Localized effect
- Dynamic scheduler keeps all processors busy

# Exploring Parallel Computation Models



## Map/Reduce Provides Coarse-Grained Parallelism

- Computation done by independent processes
- File-based communication

## Observations

- Relatively “natural” programming model
- Research issue to explore full potential and limits

# Beyond Map/Reduce

## Typical Map/Reduce Applications

- Sequence of steps, each requiring map & reduce
- Series of data transformations
- Iterating until reach convergence

## Strengths of Map/Reduce

- User writes simple functions, system manages complexities of mapping, synchronization, fault tolerance
- Very general
- Good for large-scale data analysis

## Limitations

- No locality of data or activity
- Each map/reduce step must complete before next begins

# Generalizing Map/Reduce


- Microsoft Dryad Project

## Computational Model

- Acyclic graph of operators
  - But expressed as textual program
- Each takes collection of objects and produces objects
  - Purely functional model

## Implementation Concepts

- Objects stored in files or memory
- Any object may be lost; any operator may fail
- Replicate & recompute for fault tolerance
- Dynamic scheduling
  - # Operators >> # Processors




# CMU GraphLab

- Carlos Guestrin, et al.
- Graph algorithms used in machine learning

## View Computation as Localized Updates on Graph

- New value depends on own value + those of neighbors
- Update repeatedly until converge





# Machine Learning Example

## PageRank Computation

■ Larry Page & Sergey Brinn, 1998

## Rank “Importance” of Web Pages



# PageRank Computation

## Initially


- Assign weight 1.0 to each page

## Iteratively

- Select arbitrary node and update its value

## Convergence

- Results unique, regardless of selection ordering



$$R_1 \leftarrow 0.1 + 0.9 * (\frac{1}{2} R_2 + \frac{1}{4} R_3 + \frac{1}{3} R_5)$$

# PageRank with Map/Reduce

$$R_1 \leftarrow 0.1 + 0.9 * (\frac{1}{2} R_2 + \frac{1}{4} R_3 + \frac{1}{3} R_5)$$

## Each Iteration: Update all nodes

- **Map:** Generate values to pass along each edge
  - Key value 1:  $(1, \frac{1}{2} R_2)$      $(1, \frac{1}{4} R_3)$                  $(1, \frac{1}{3} R_5)$
  - Similar for all other keys
- **Reduce:** Combine edge values to get new rank
  - $R_1 \leftarrow 0.1 + 0.9 * (\frac{1}{2} R_2 + \frac{1}{4} R_3 + \frac{1}{3} R_5)$
  - Similar for all other nodes

## Performance

- Very slow!
- Altavista Webgraph 2002
  - 1.4B vertices, 6.7B edges

Hadoop	800 cores	9000s
--------	-----------	-------

# PageRank with GraphLab

## Operation

- **Graph partitioned across multiple processors**

- Each doing updates to its portion of graph
- Exploits locality
- Greater asynchrony
- Only iterate over portions of graph where values are changing

## Performance

- **Altavista Webgraph 2002**

- 1.4B vertices, 6.7B edges

<b>Hadoop</b>	<b>800 cores</b>	<b>9000s</b>
<b>Prototype GraphLab2</b>	<b>512 cores</b>	<b>431s</b>

# Conclusions

## Distributed Systems Concepts Lead to Scalable Machines

- Loosely coupled execution model
- Lowers cost of procurement & operation

## Map/Reduce Gaining Widespread Use

- Hadoop makes it widely available
- Great for some applications, good enough for many others

## Lots of Work to be Done

- Richer set of programming models and implementations
- Expanding range of applicability
  - Problems that are data *and* compute intensive
  - The future of supercomputing?